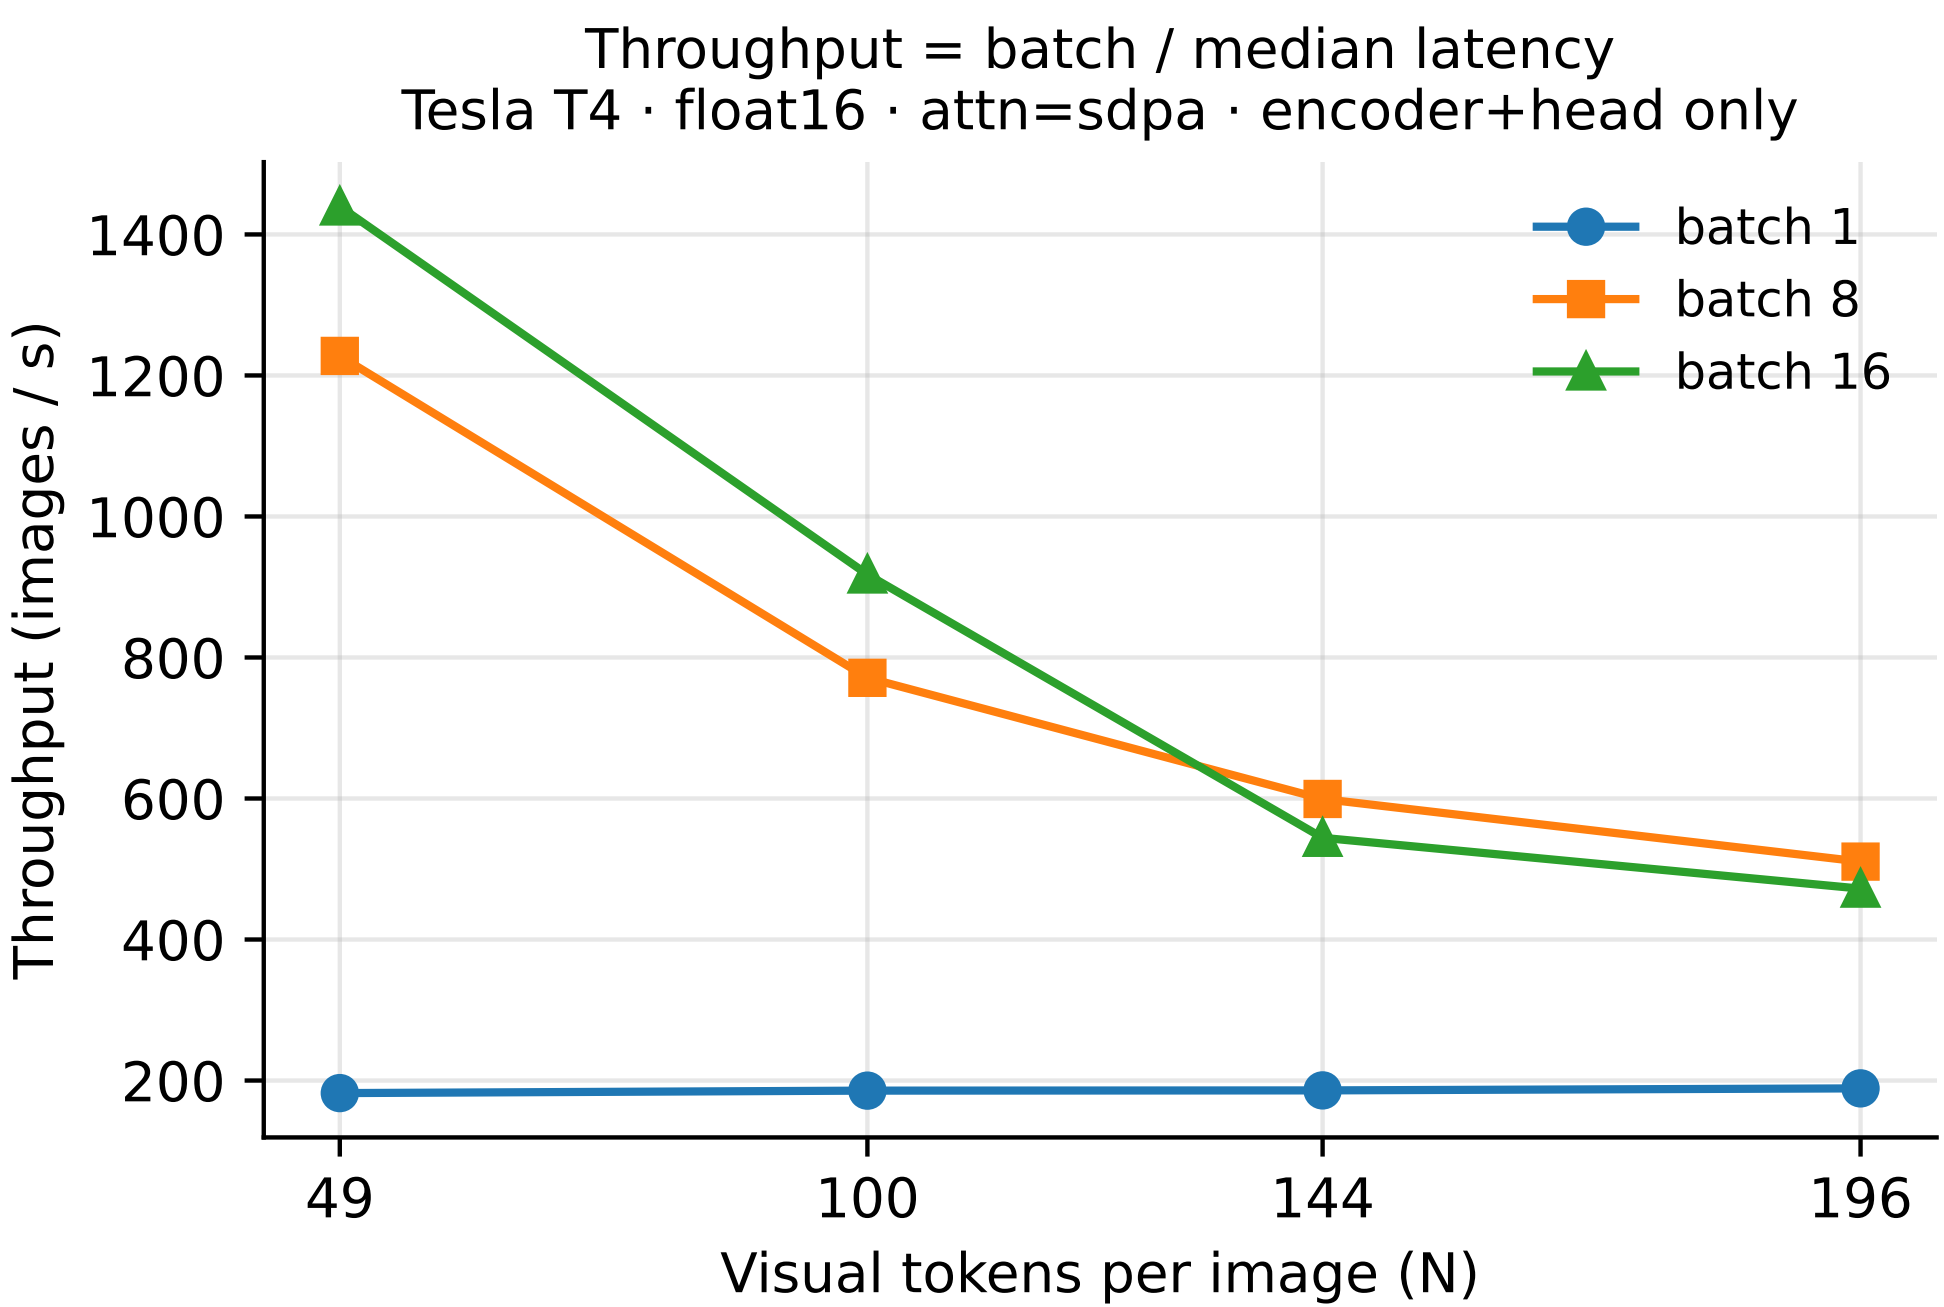
Throughput = batch / median latency
Tesla T4 · float16 · attn=sdpa · encoder+head only
batch 1
batch 8
batch 16
Throughput (images / s)
1400
1200
1000
800
600
400
200
49
100
144
196
Visual tokens per image (N)