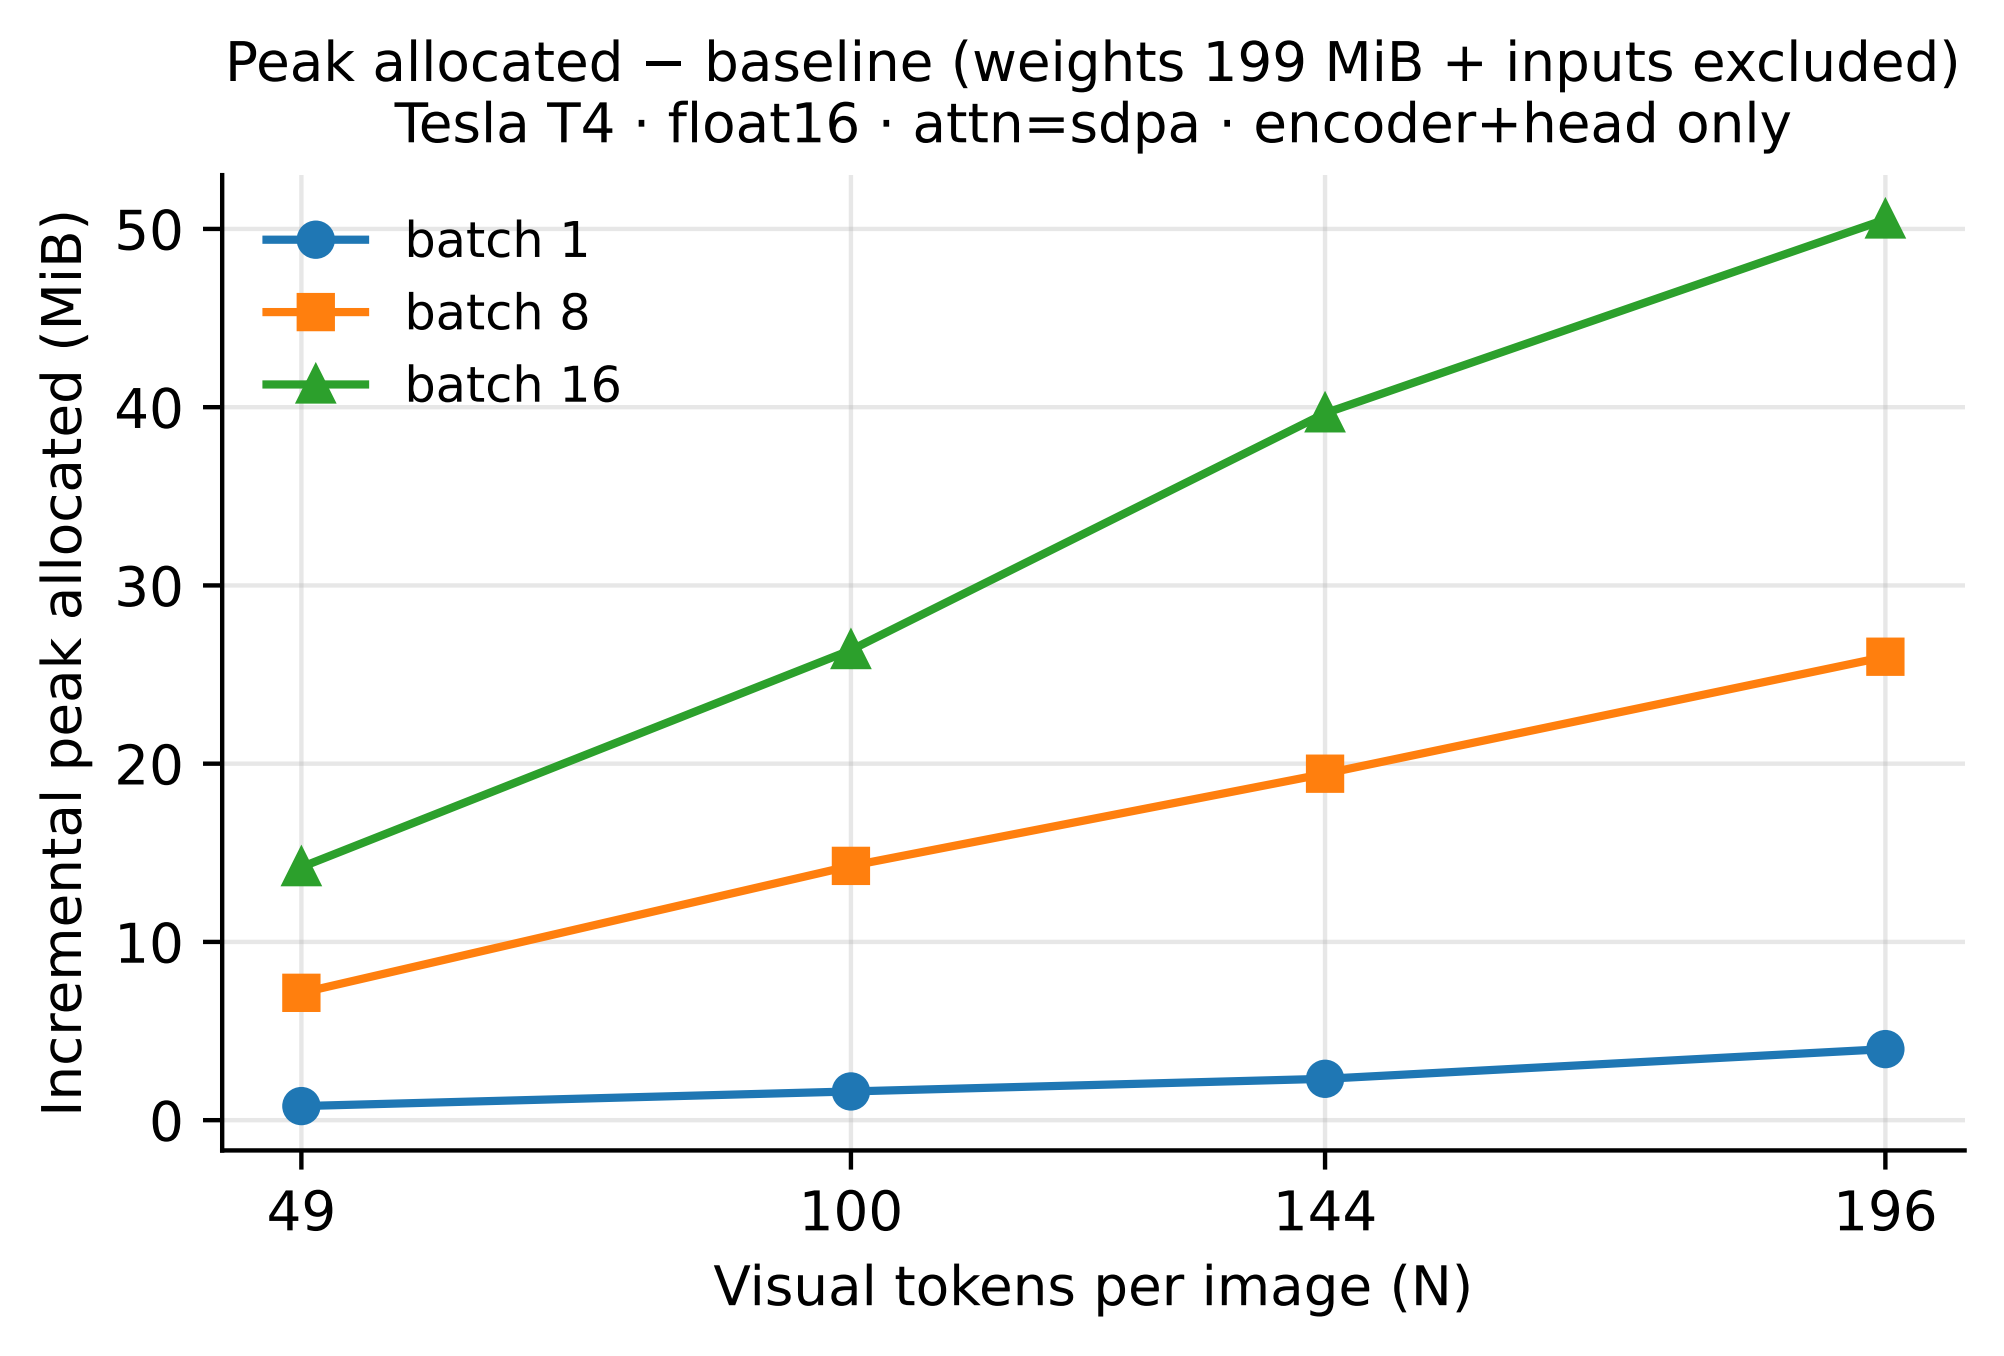
Peak allocated − baseline (weights 199 MiB + inputs excluded)
Tesla T4 · float16 · attn=sdpa · encoder+head only
batch 1
batch 8
batch 16
Incremental peak allocated (MiB)
0
10
20
30
40
50
49
100
144
196
Visual tokens per image (N)