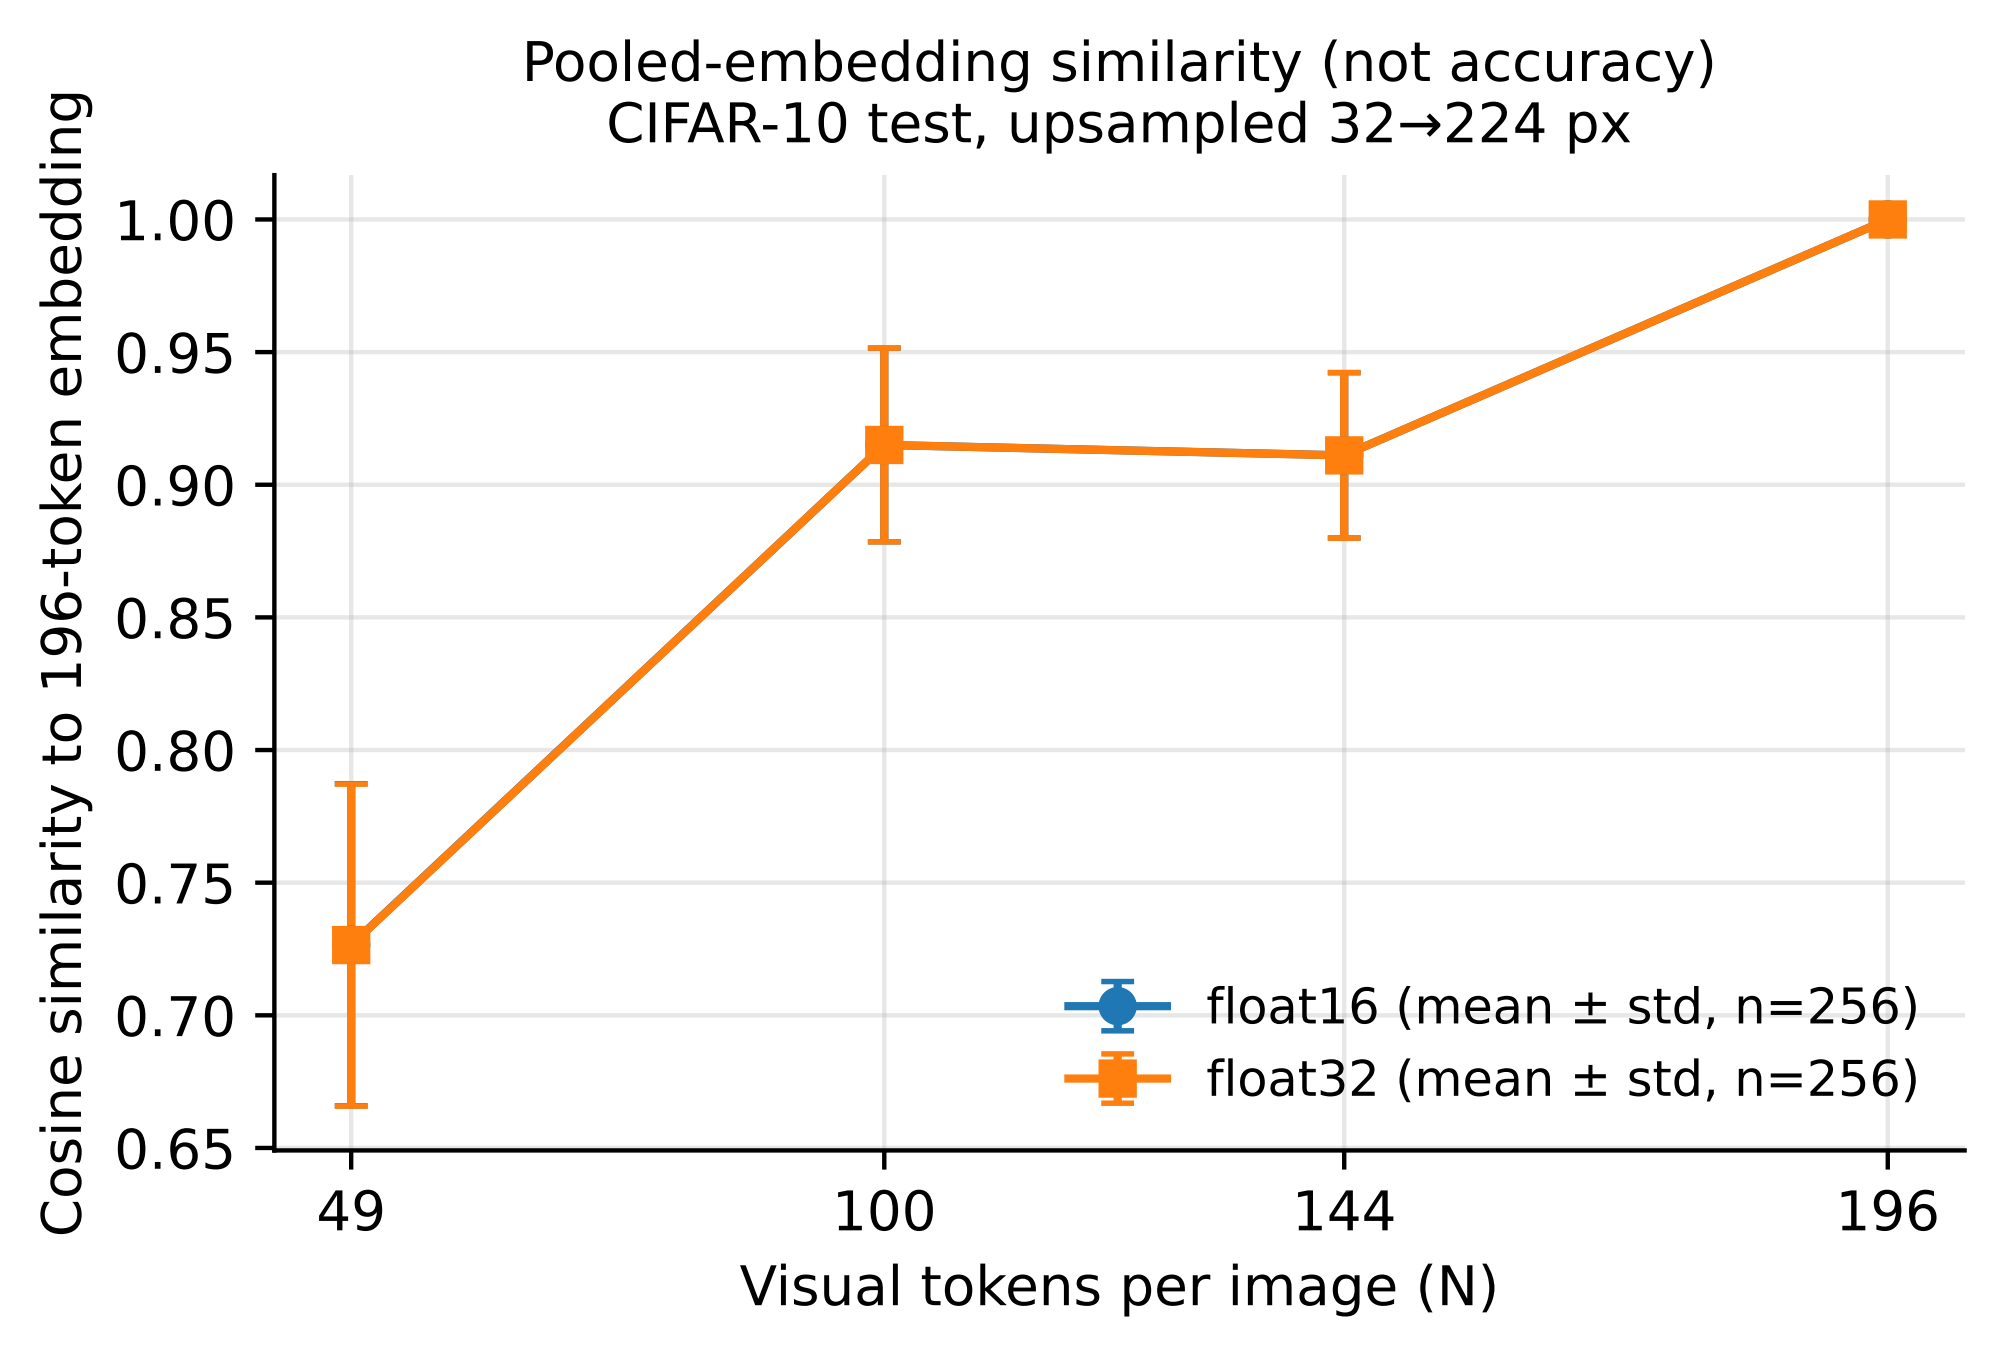
Pooled-embedding similarity (not accuracy)
CIFAR-10 test, upsampled 32→224 px
Cosine similarity to 196-token embedding
1.00
0.95
0.90
0.85
0.80
0.75
0.70
0.65
49
100
144
196
Visual tokens per image (N)
float16 (mean ± std, n=256)
float32 (mean ± std, n=256)