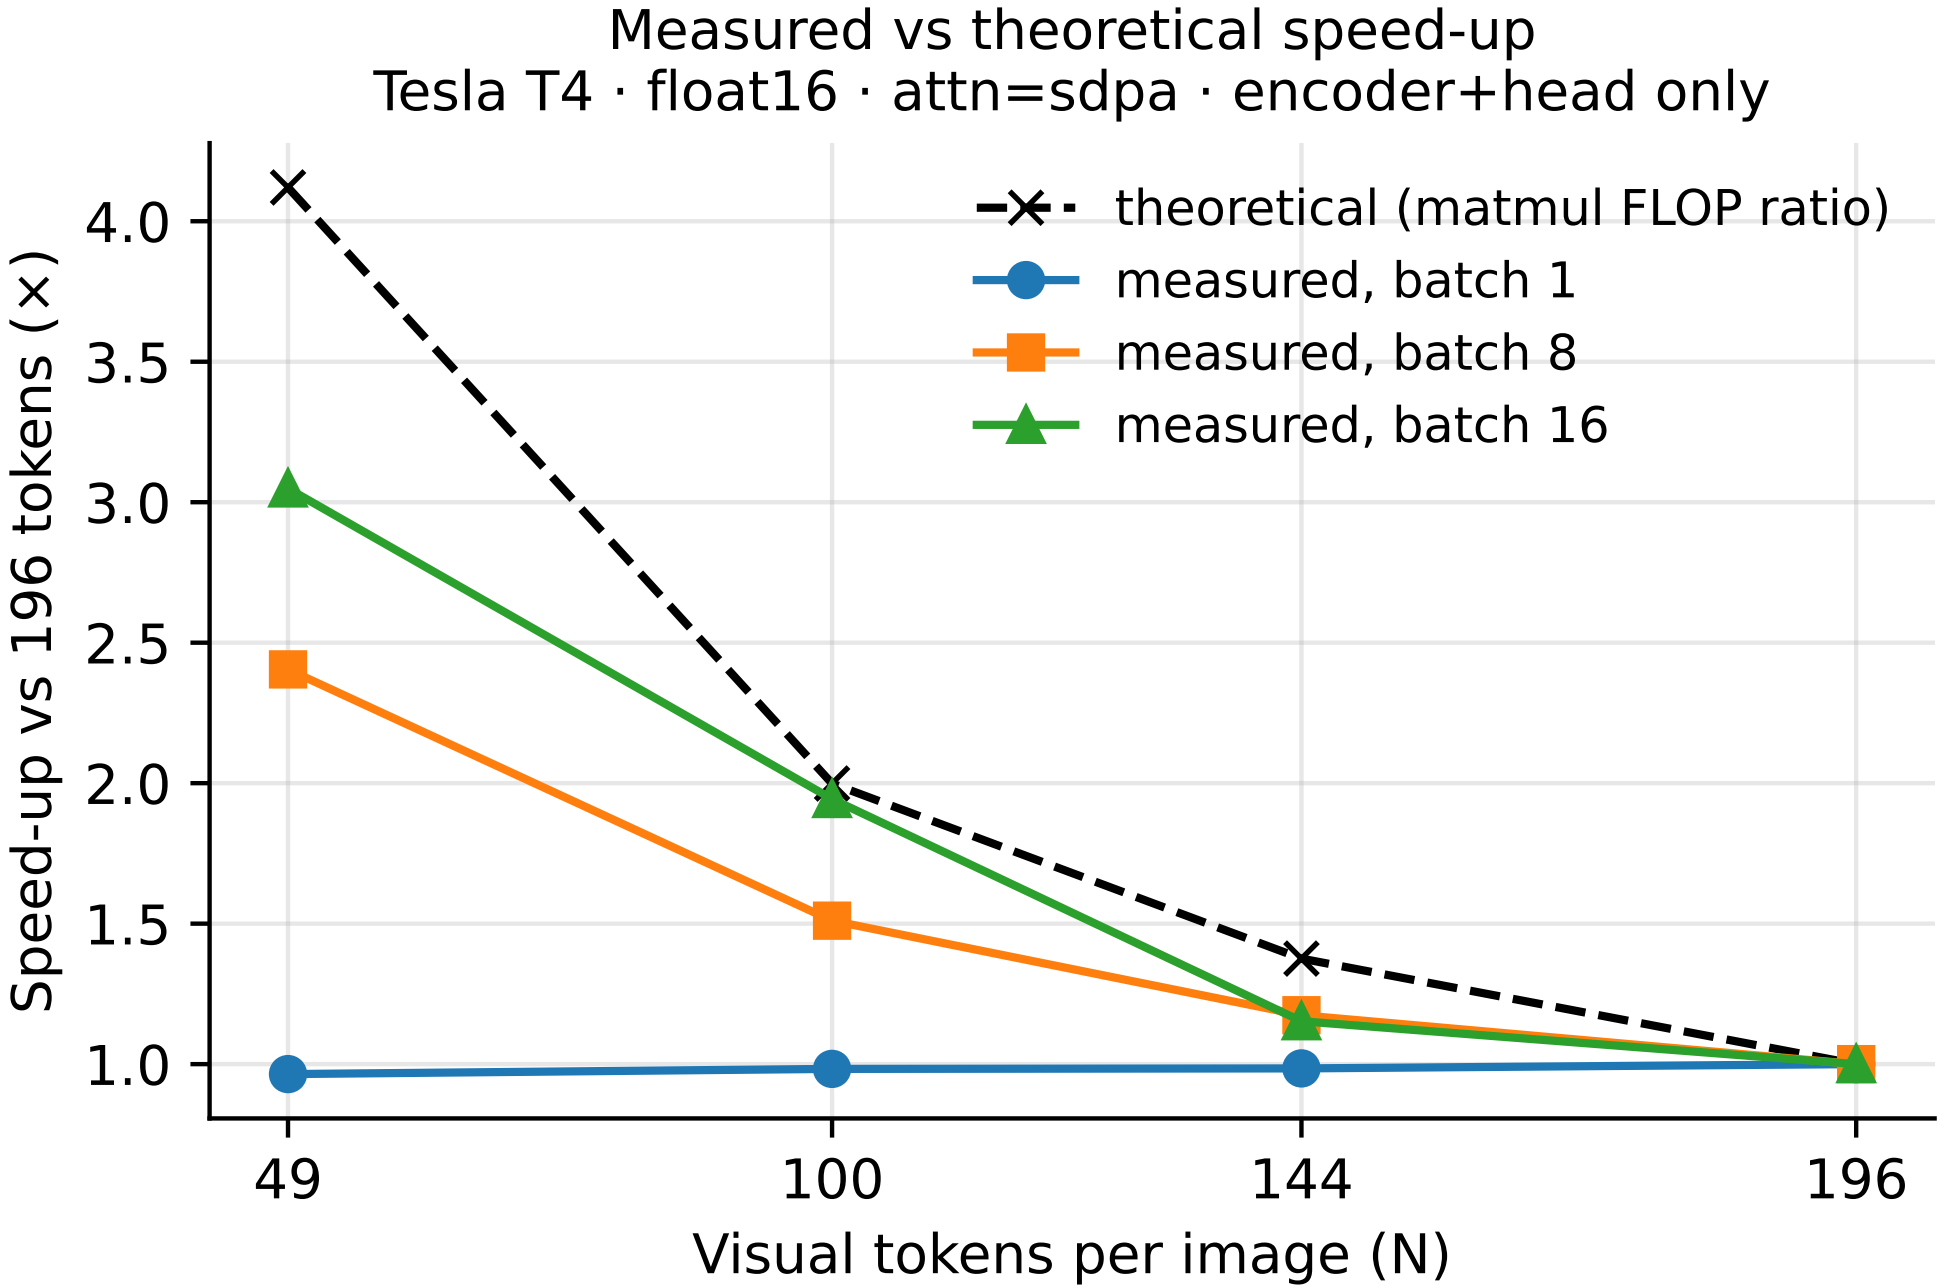
Measured vs theoretical speed-up
Tesla T4 · float16 · attn=sdpa · encoder+head only
theoretical (matmul FLOP ratio)
measured, batch 1
measured, batch 8
measured, batch 16
Speed-up vs 196 tokens (×)
4.0
3.5
3.0
2.5
2.0
1.5
1.0
49
100
144
196
Visual tokens per image (N)